



Increasing your performance

MySQL Enterprise Thread Pooling

Dale Dasker

Manager, MySQL Solution Engineering

dale.dasker@oracle.com

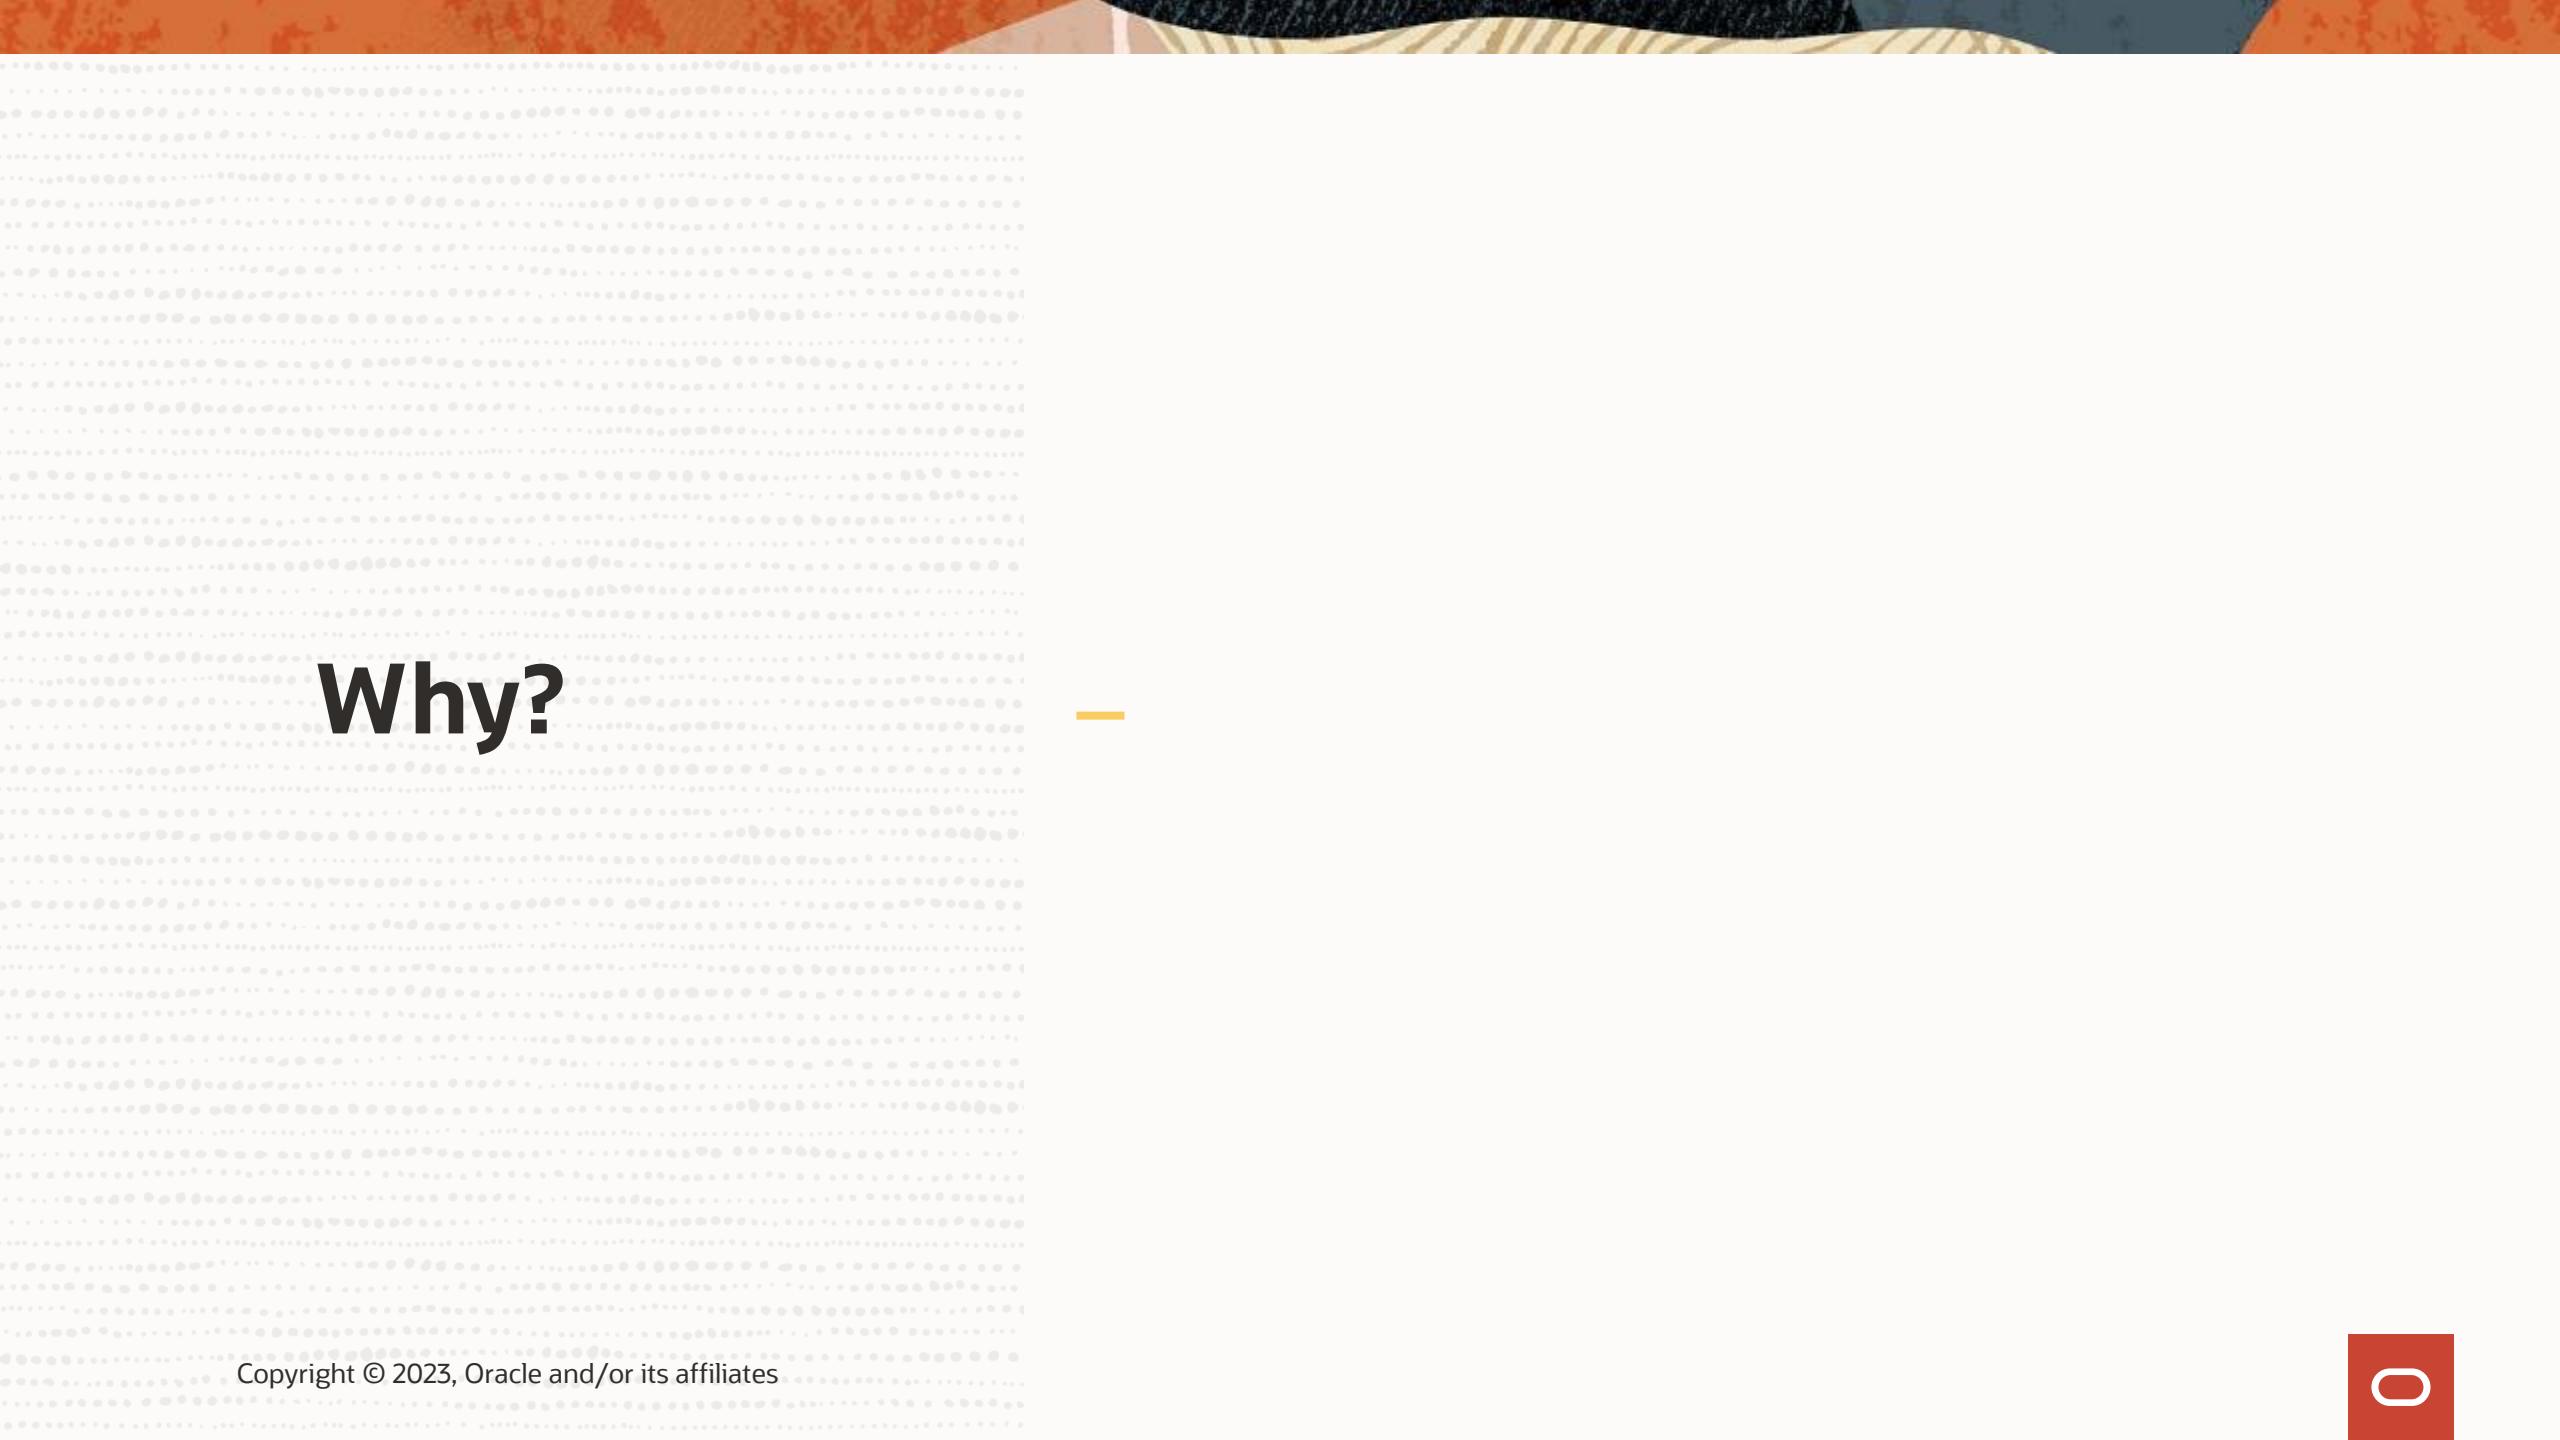
August 24, 2023

Agenda

Achieve Increased Throughput and Scalability with Enterprise Thread Pooling

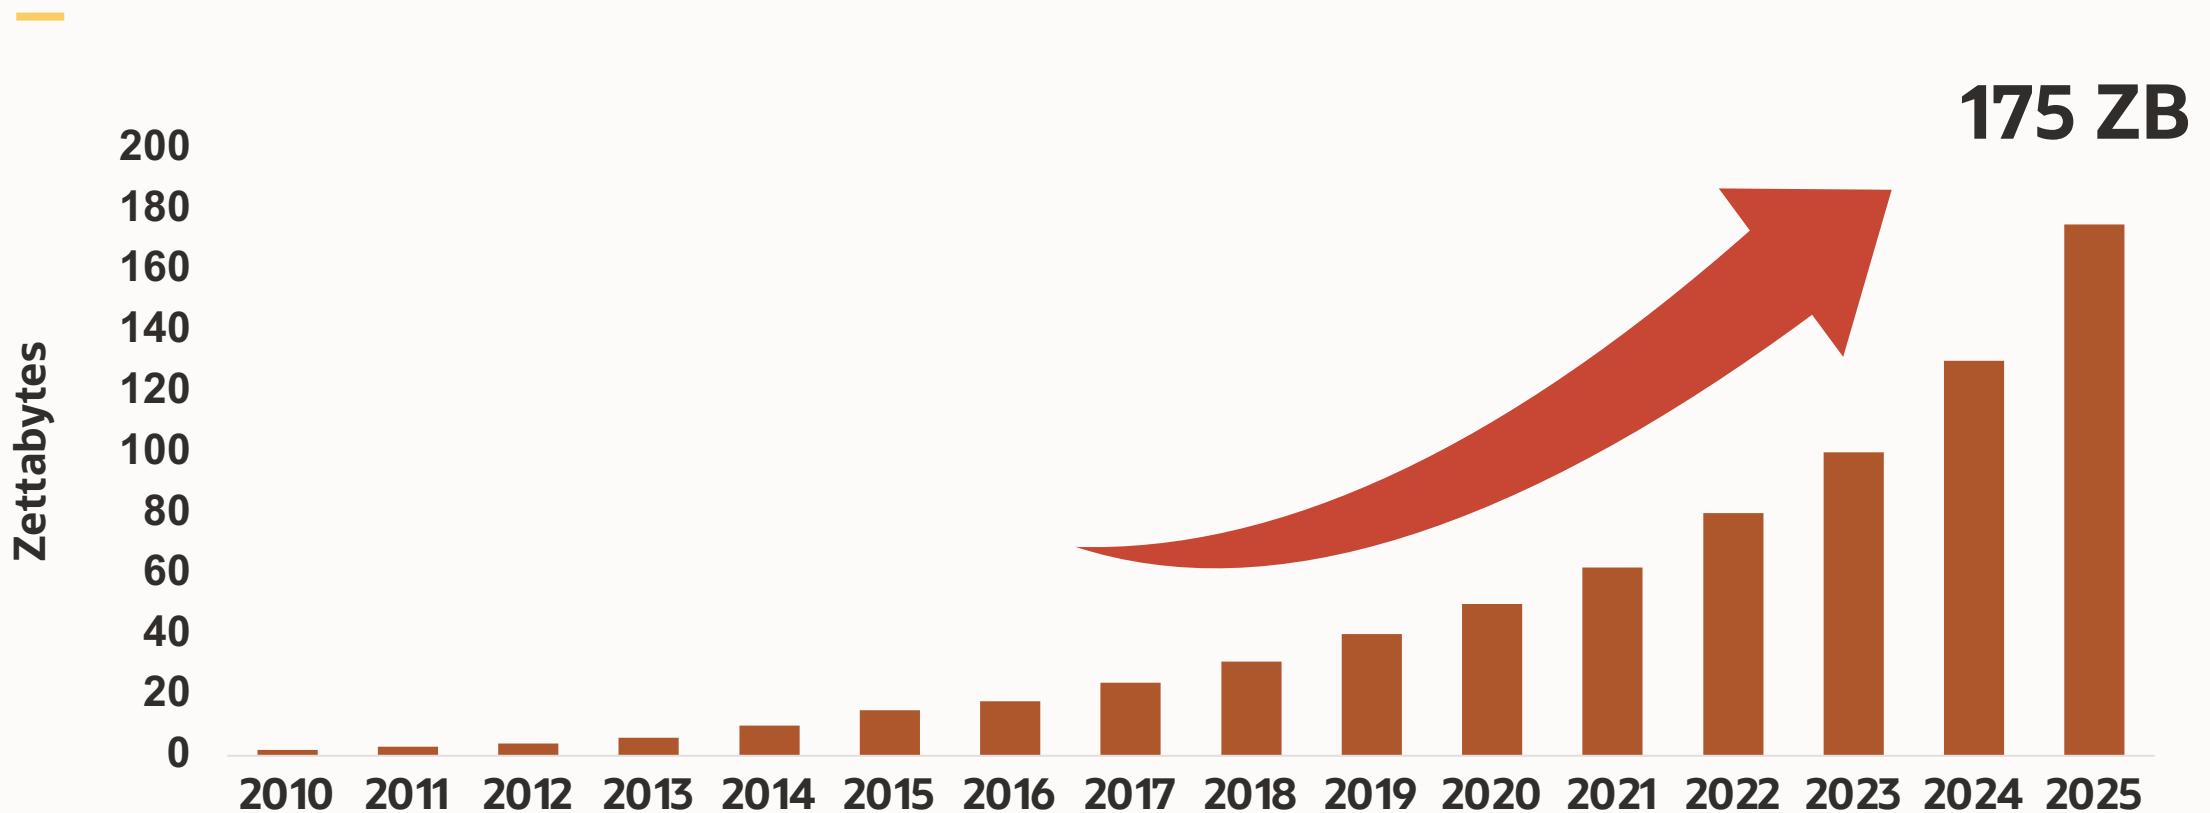
- Thread Pooling Overview
- Sysbench Overview
- Setup and Installation of:
 - Enterprise Thread Pool
 - sysBench

Why?



—

Global Datasphere



MySQL Version / CPU Support Evolution

2008

2009

2010

2012

2015

2018

MySQL 5.0
Up to 4 Cores

MySQL 5.1
(InnoDB Plugin)
Up to 16 Cores
(Sun Micro)

MySQL 5.5
Up to 32 Cores
(Oracle)

MySQL 5.6
Up to 48 Cores

MySQL 5.7
72 Cores

MySQL 8.0
96 Cores

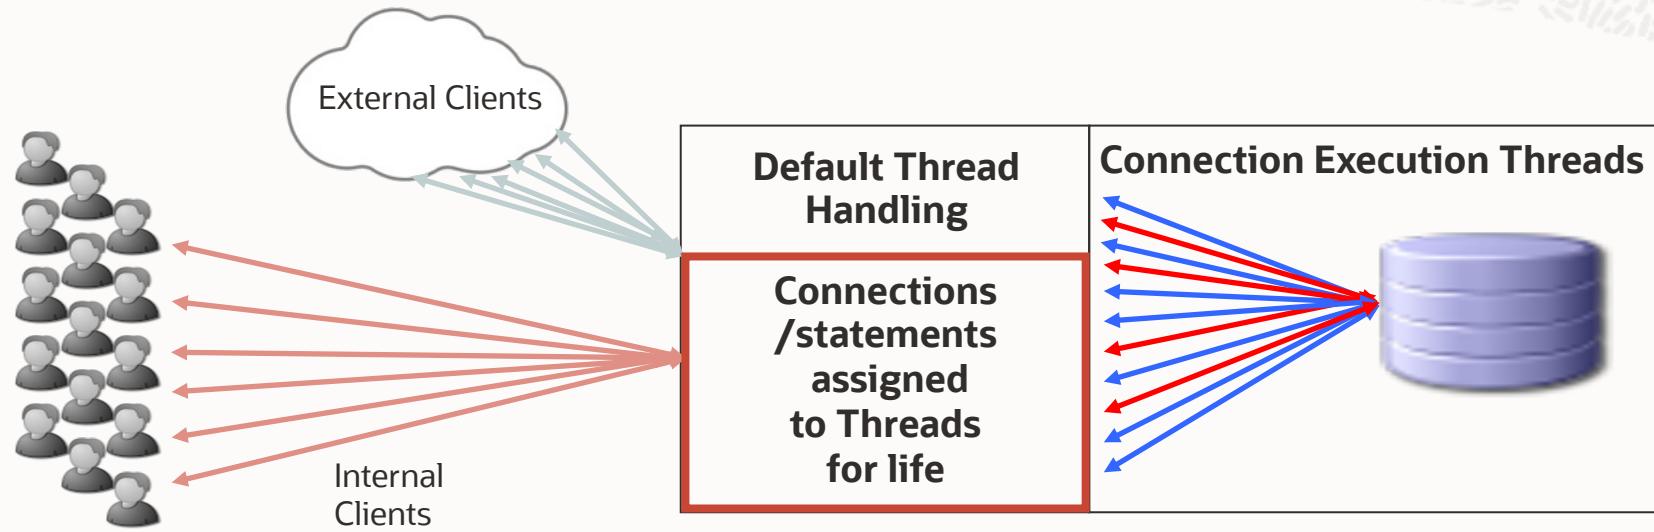


Thread Pool: Problem Definition

- Current mapping from connection to thread is one thread per connection.
 - Memory usage increases
 - More Cache Misses
 - Higher contention on “hot spots”
- Write-Intensive workloads more affected

MySQL Enterprise Thread Pool

Default Connection Behavior



- Connections assigned to 1 thread for the life of the connection, same thread used for all statements
- No prioritization of threads, statement executions
- Many concurrent connections = many concurrent execution threads to consume server memory, limit scalability

Thread Pool Solution

- Wait for execution of a query until the MySQL Server has sufficient CPU and memory resources to execute it.
- Prioritize queries on connections that have an ongoing transaction.
- Split threads into groups individually handled to avoid making the solution a problem in itself, aim is to manage one active thread per group.
- Avoid deadlocks when queries are stalled or execute for a long time.

Thread Pool Solution: Avoiding Scalability Issue

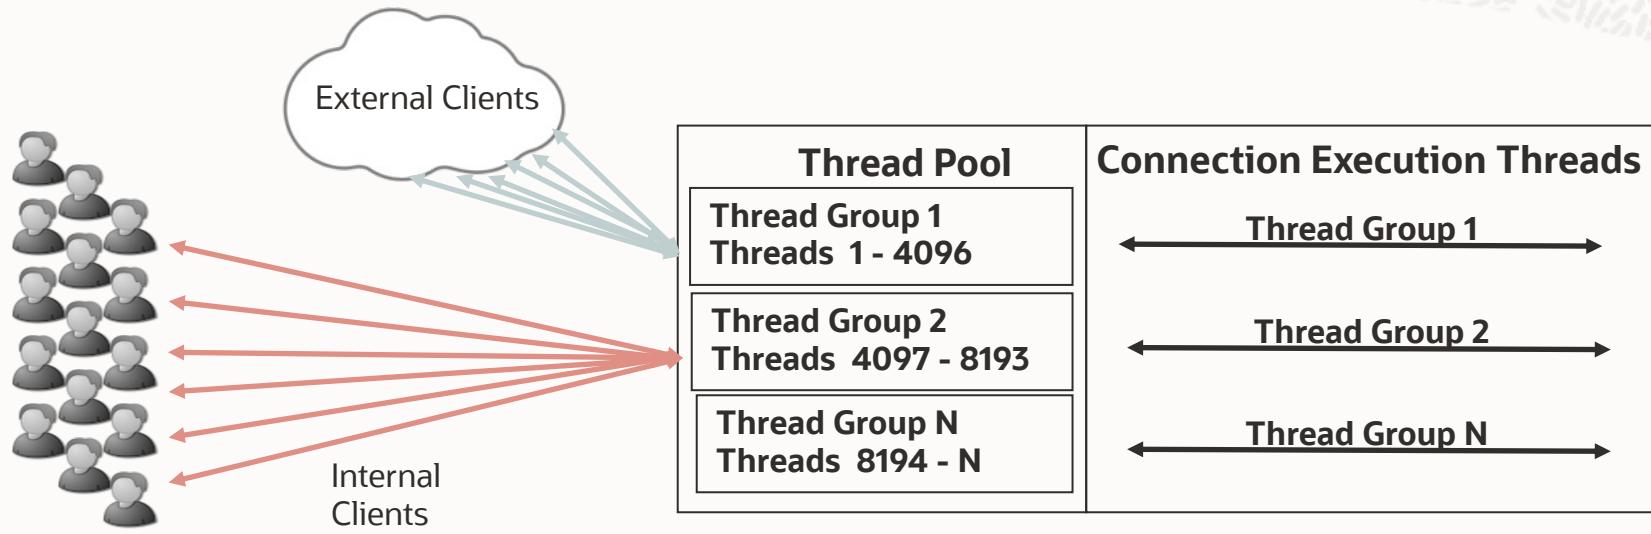
- In implementing a thread pool it is easy to cause a scalability problem by managing all threads in one pool.
- To avoid this issue we have divided threads and connections into **thread groups**. The aim of the implementation is to have one or zero queries executing per thread group at any time.
- Connections are bound to thread groups by round robin assignment at connect time.
- Each thread group maintains prioritized queues of connections awaiting execution.
- So problem solved using Divide-and-Conquer technique.

Thread Pool Solution: Minimize number of concurrent transactions

- To avoid scalability bottlenecks due to too many concurrent transactions it is imperative to limit the number of concurrent transactions.
- This is accomplished by prioritizing connections waiting to execute. Connections that have ongoing transactions are queued in a high priority queue and others are queued in the low priority connection.
- This means that low priority connections will wait until all high priority connections have completed.

MySQL Enterprise Thread Pool

Scales user connections



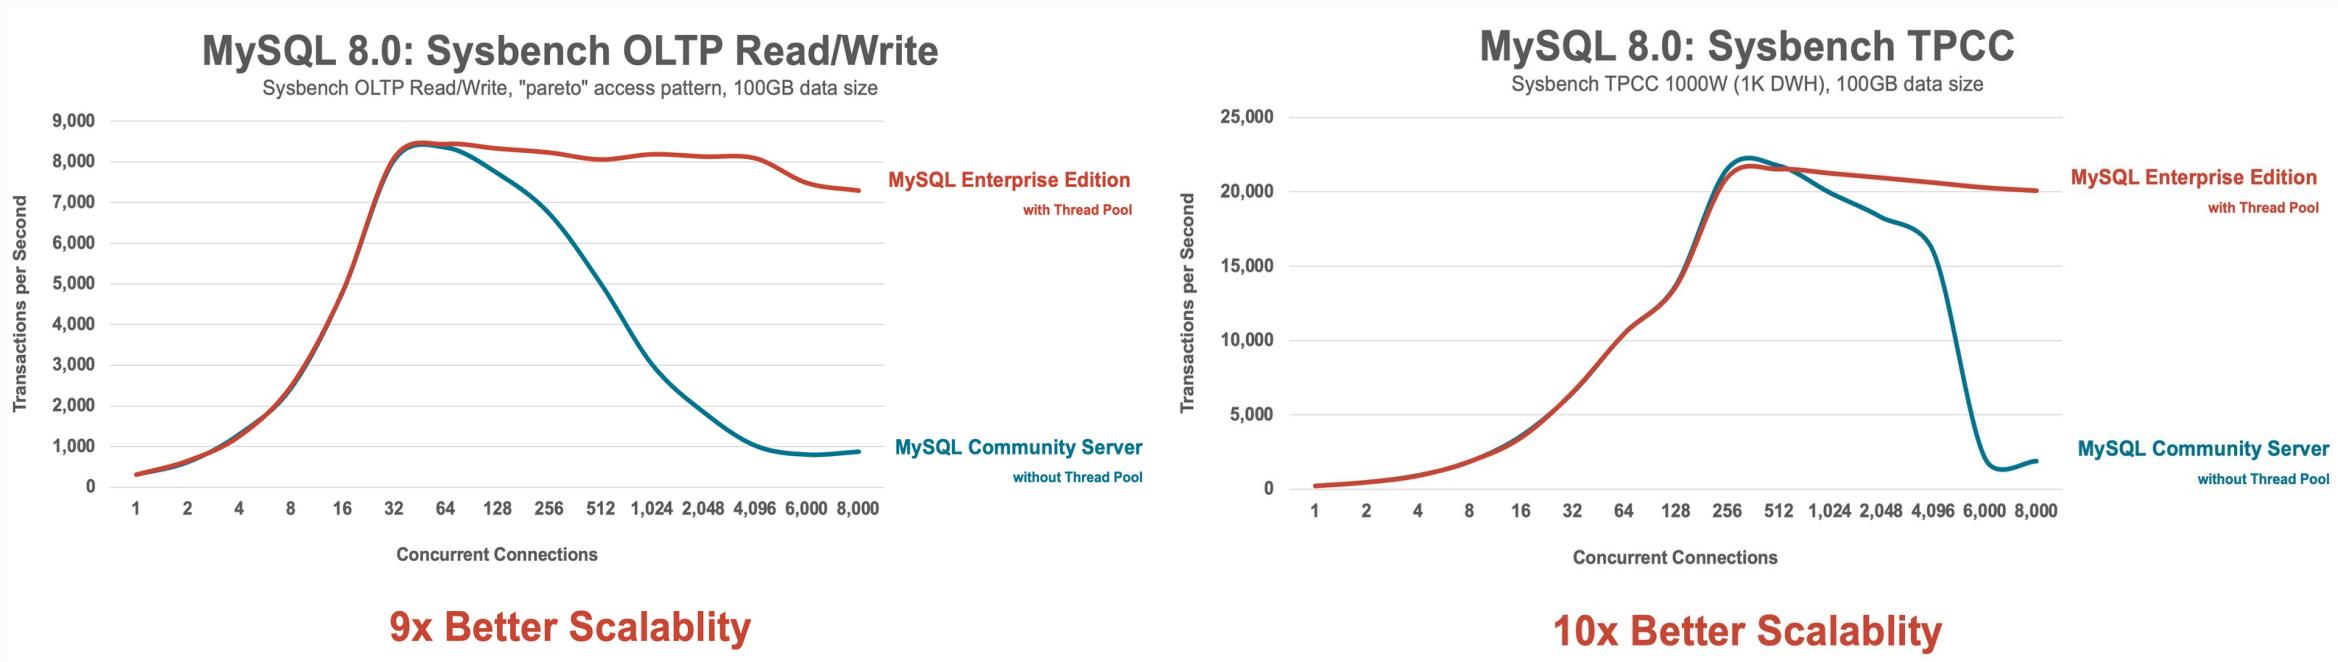
Provided by the thread pool plugin

More intelligent model than the default

- Improves scalability as concurrent connections grow
- Threads are prioritized and statements queued

Performance Scalability with MySQL Enterprise Thread Pool

- To meet the sustained performance and scalability of ever increasing user
- A highly scalable thread-handling model designed to reduce overhead in managing client connections and statement execution threads



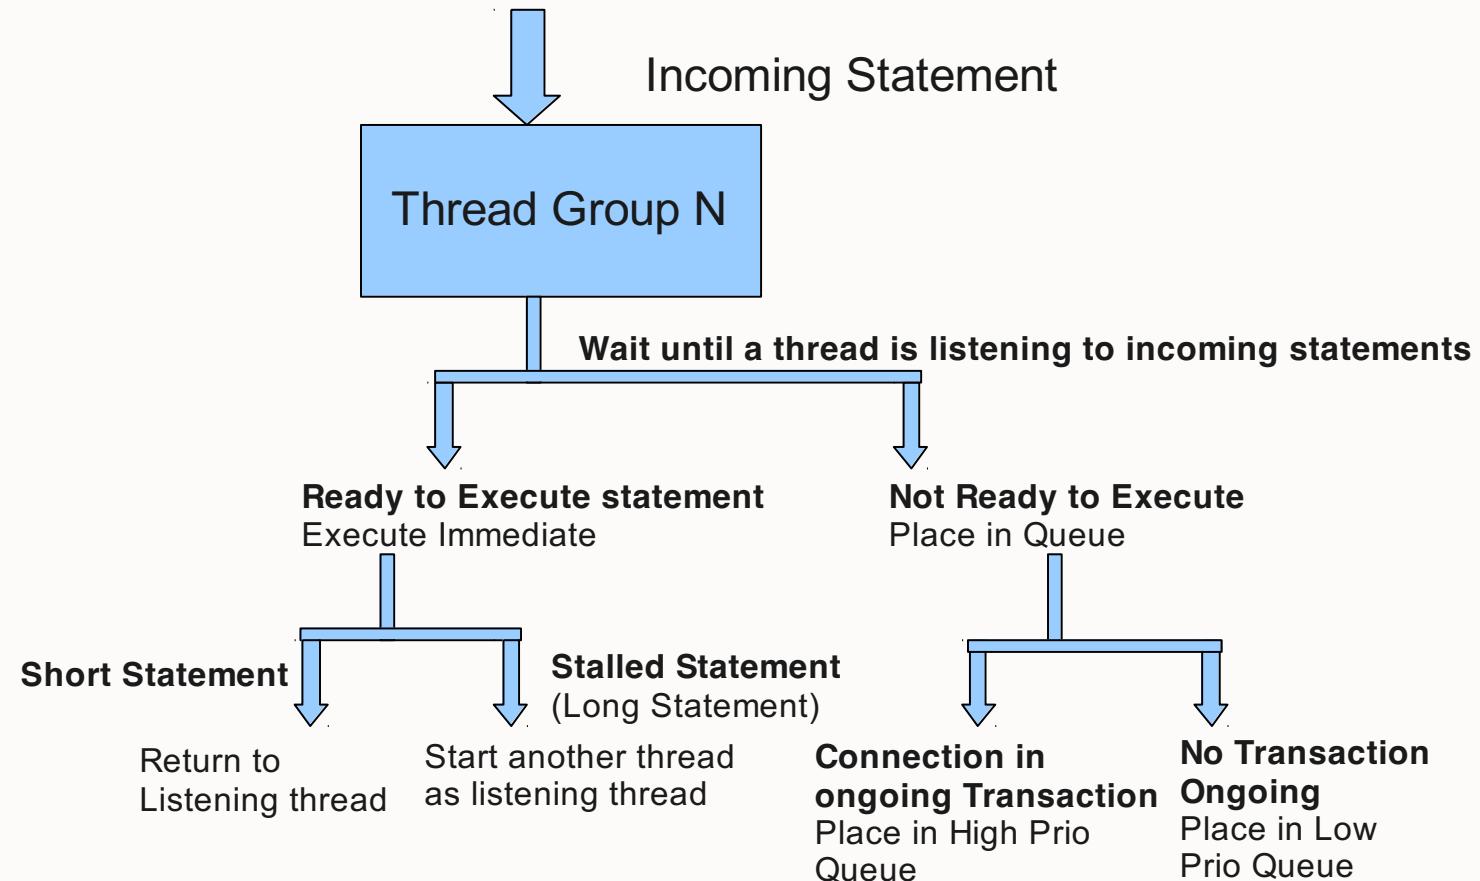
Server: 2 Sockets, 48cores, 96 threads Intel(R) Xeon(R) Platinum 8268 CPU @ 2.90GHz, RAM: 192GB
Storage: NVMe Optane 2 x 375GB, OS: OL7.9 UEK6, MySQL: v8.0.31 OpenSSL-1.1.1

Client: 2 Sockets, 44cores, 88 threads Intel(R) Xeon(R) CPU E5-2699 v4 @ 2.20GHz

MySQL Enterprise Thread Pool

Flow Logic

Thread Pool Operations



Thread Pool Tuning

Minimize number of concurrent transactions

- The prioritization of long transactions can have negative effect on short transactions, even causing a live lock situation.
- To ensure that there is a remedy for this problem we move queries from low priority queue to high priority queue when they have waited for a configurable time.
- This time is configurable in the parameter **--thread_pool_prio_kickup_timer**
- We will never move more than one query per 10 ms per thread group.

Thread Pool Tuning

Limit number of concurrent queries executed

- The thread pool aims to limit the number of concurrent queries within each thread group to zero or one.
- However there is a live lock issue that needs to be handled. A long-running query could potentially block all other queries in its thread group forever (the replication connection has some “eternal queries”).
- To avoid this a query that executes for a long time is declared as stalled, the time until a query is stalled is controlled through the configuration parameter --
thread_pool_stall_timer.
- When a query is declared as stalled, another query is allowed to start in this thread group.

Thread Pool vs. Connection Pool

- Given that a connection has been equal to a thread in MySQL, it is easy to confuse the Connection Pool with the Thread Pool.
- They are however solutions to different problems. A connection pool makes it possible to reuse connections and thus saving on having to execute connects to the MySQL Server.
- The thread pool works on the server side to limit the number of concurrently executing queries.
- Thus they can be used independently of each other.

Thread Pool: When to use?

- The most important variable to monitor is `threads_running`.
- This variable provides information about number of concurrent queries executed in MySQL Server.
- It is sufficient if this variable has spikes that put it above the optimal `thread_pool_size` for the type of queries used in the MySQL Server for the thread pool to be useful.
- Given that these spikes usually happen in overload situations, the thread pool is an important protection against troubles in overload situations.

https://dev.mysql.com/doc/refman/8.0/en/server-status-variables.html#statvar_Threads_running

https://dev.mysql.com/doc/refman/8.0/en/server-system-variables.html#sysvar_thread_pool_size

Thread Pool: When to use (cont.)?

- If you are currently using **-innodb-thread-concurrency** to limit number of concurrent queries, then usage of the thread pool is very likely to be beneficial.
- Also if most of your queries are short queries then the thread pool is likely to provide benefits.
- Long queries are not negative for the thread pool since it will operate very similarly to the normal MySQL threading model for these queries. In these cases there will be fewer benefits to enabling the thread pool.

Thread Pool: Operations

- **thread_pool_dedicated_listeners**: Dedicates a listener thread in each thread group to listen for incoming statements from connections assigned to the group.
- **thread_pool_max_transactions_limit**: The maximum number of transactions permitted by the thread pool plugin.
- **thread_pool_query_threads_per_group**: The number of query threads permitted in a thread group (the default is a single query thread).
- **thread_pool_size**: The number of thread groups in the thread pool. This is the most important parameter controlling thread pool performance.
- **thread_pool_stall_limit**: The time before an executing statement is considered to be stalled.
- **thread_pool_transaction_delay**: The delay period before starting a new transaction.
- **thread_pool_algorithm**: The concurrency algorithm to use for scheduling.
- **thread_pool_high_priority_connection**: How to schedule statement execution for a session.
- **thread_pool_max_active_query_threads**: How many active threads per group to permit.
- **thread_pool_max_unused_threads**: How many sleeping threads to permit.
- **thread_pool_prio_kickup_timer**: How long before the thread pool moves a statement awaiting execution from the low-priority queue to the high-priority queue.

Thread Pool: Optimal configuration

--thread_pool_size (Number of thread groups)

- InnoDB Read Workloads usually around 30-40
- InnoDB Write Workloads lower at around 12-30
- InnoDB Read/Write **Workloads at 16-36**
- Most important setting. Only set at server startup.
- Also is dependent on type of queries in workload, but default setting of 16 is usually a good starting point.
- Beneficial often to lock MySQL Server to same number of CPU threads as this parameter is set to.

Thread Pool: Optimal configuration (cont.)

--thread_pool_query_threads_per_group (Number of query thread in thread groups)

- Regular workloads – 2
- Also is dependent on type of queries in workload, but default setting of 1 is usually a good starting point.
- Beneficial often to lock MySQL Server to same number of CPU threads as this parameter is set to.

https://dev.mysql.com/doc/refman/8.0/en/server-system-variables.html#sysvar_thread_pool_query_threads_per_group

Thread Pool: Optimal configuration (cont.)

--thread_pool_stall_limit

- For benchmarks usually set to at least 1 second (=100).
- Check thread pool Information/Performance Schema tables to ensure not too many queries are stalled.
- Default setting is conservative to avoid impact of many long queries (=6, 60ms).
- Measured in 10ms segments. Setting of 6 = 60ms.
- Dynamic variable. Can be altered during runtime.
- Determine fraction of queries that are stalled (lower is better):
 - ```
SELECT SUM(STALLED_QUERIES_EXECUTED) / SUM(QUERIES_EXECUTED) FROM performance_schema.tp_thread_group_stats;
```

## Thread Pool: Optimal configuration (cont.)

---

### **--thread\_pool\_prio\_kickup\_timer**

- For benchmarks usually set to at least 10 seconds (=10000).
- Set according to risk of livelock due to long-running queries.
- Default setting of 1 second should be OK for most cases, increasing this value is ok if the environment has only few and controlled long queries.

# Thread Pool

## Information Schema Tables

---

3 Tables:

### **TP\_THREAD\_STATE**

- Shows current state of all threads controlled by thread pool

### **TP\_THREAD\_GROUP\_STATE**

- Shows current state of all thread groups in the thread pool

### **TP\_THREAD\_GROUP\_STATS**

- Shows statistics for each of the thread groups in the thread pool

# Thread Pool Information Schema Tables:

## TP\_THREAD\_STATE

---

Can view

- how long time a query executed in thread
- and if in a waiting state from where wait was initiated

<https://dev.mysql.com/doc/refman/8.0/en/performance-schema-tp-thread-state-table.html>

# Thread Pool Information Schema Tables:

## TP\_THREAD\_GROUP\_STATE

---

Can view current values for example:

- Number of threads in thread group
- Number of connections in the queues
- Number of threads in various categories
- Number of connections per thread group
- Configuration settings
- For any listening thread
- Query with current longest wait in queues
- Number of active and stalled queries

<https://dev.mysql.com/doc/refman/8.0/en/performance-schema-tp-thread-group-state-table.html>

# Thread Pool Information Schema Tables:

## TP\_THREAD\_GROUP\_STATS

---

Can view:

- Thread group ID
- Number of connections started, closed
- Number of statements executed, queued
- Number of threads started
- Stats on consumer, reserve, waiter thread roles
- Stats on waits and locking

<https://dev.mysql.com/doc/refman/8.0/en/performance-schema-tp-thread-group-stats-table.html>

# Thread Pool: Summary

Problem and how it is solved

---



## Problem 1:

How to split threads into groups individually handled to avoid making the solution a problem in itself, aim is to manage one active thread per group?

## Solution:

Connections are put into a thread group at connect time by round robin. Configurable amount of thread groups. This ensures that the thread pool itself isn't a scalability hog.

# Thread Pool: Summary

Problem and how it is solved

---

## Problem 2:

How to wait for execution of a query until the MySQL Server has sufficient CPU and memory resources to execute it?

## Solution:

Each thread group tries to keep the number of executing queries to one or zero. If a query is already executing in the thread group, put connection in wait queue.

# Thread Pool: Summary

Problem and how it is solved

---

## Problem 3:

How to prioritize queries on connections that have an ongoing transaction?

## Solution:

Put waiting connections in high priority queue when a transaction is already started on the connection.

# Thread Pool: Summary

Problem and how it is solved

---

## Problem 4:

How to avoid deadlocks when queries are stalled or execute for a long time?

## Solution:

Allow another query to execute when the executing query in the thread group is declared as stalled (after a configurable time).

# Benchmarking

---

# Popular Generic Workloads

---

## Sysbench

- Our recommendation
- Looks simple, but still the most complete test kit !
- OLTP, RO/RW, points on various RO and RW issues

## TPC-C :

- DBT-2 / TPCC-like / HammerDB / Sysbench-TPCC

## TPC-H / DBT-3

- DWH, RO, complex heavy queries, loved by Optimizer Team ;-)

## dbSTRESS

- OLTP, RO/RW, several tables, points on RW and Sec.IDX issues

# Why Sysbench is Favorite Choice

---

- Simple to install & use
- Covers most important MySQL Workloads
- Very flexible
  - LUA scripts
  - Easy to write your own tests
- Source available
  - <https://github.com/akopytov/sysbench>

## Why Sysbench is Favorite Choice (cont.)

---

- Stress hardware
  - CPU
  - Fileio
  - Memory
  - Threads
- Compare Cloud nodes
  - Queries per \$
- Assess feasibility of designs
  - Group Replication across broad networks
- Great for synthetic benchmarks

# Installation & Usage (RPM)

---



Add the EPEL library:

- sudo yum -y install <https://dl.fedoraproject.org/pub/epel/epel-release-latest-8.noarch.rpm>

Install Sysbench:

- sudo yum -y install sysbench

Verify installation:

- sysbench --version

Useful Links:

- <https://github.com/akopytov/sysbench>
- <https://www.mortensi.com/2021/06/how-to-install-and-use-sysbench/>

# Installation & Usage (RPM) cont.



## Simple test of cpu:

- sysbench cpu --cpu-max-prime=10000 run

## Simple test of MySQL:

- Create database schema
  - mysql -uroot -pWelcome1! -e"CREATE DATABASE sysbench"
- Prepare databases
  - sysbench oltp\_read\_write --table-size=1000000 --db-driver=mysql --mysql-host=127.0.0.1 --mysql-db=sysbench --mysql-user=root --mysql-password>Welcome1!  
**Prepare**
- Run sysbench
  - sysbench oltp\_read\_write --threads=2 --report-interval=3 --histogram --time=50 --table-size=1000000 --db-driver=mysql --mysql-host=127.0.0.1 --mysql-db=sysbench --mysql-user=root --mysql-password>Welcome1! **Run**
- Clean up
  - sysbench oltp\_read\_write --db-driver=mysql --mysql-host=127.0.0.1 --mysql-db=sysbench --mysql-user=root --mysql-password>Welcome1! **cleanup**

# BMK Toolkit

- Toolkit created by Dimitri Kravtchuk (MySQL Performance Engineer)
- Designed to be an All-in-One toolkit.
  - Contains pre-built Sysbench binary statically compiled with the MySQL Client Library.
  - Fully contained tree pointed to by BMK\_HOME environment variable
  - Includes TPCC & dbSTRESS
- Sample run:

```
$ bash /BMK/sb_exec/sb11-Prepare_50M_8tab-InnoDB.sh 32 # prepare data
$ for users in 1 2 4 8 16 32 64 128 256 512 1024
do
 # run OLTP_RW for 5min each load level..
 bash /BMK/sb_exec/sb11-OLTP_RW_50M_8tab-uniform-ps-trx.sh $users 300
 sleep 15
done
```

- Download:
  - <http://dimitrik.free.fr/BMK-kit.tgz>
- Docs
  - <http://dimitrik.free.fr/blog/posts/mysql-perf-bmk-kit.html>

# Workshop Overview

---

- **Goals:**
  1. Create a OCI Compute server for hosting MySQL Enterprise Edition
  2. Install MySQL Enterprise Edition
  3. Overview & Setup Enterprise Scalability.
- **What this Workshop is not:**
  - In-depth tutorial on Oracle Cloud Infrastructure
  - MySQL Training Class
- **Lab:**
  - <https://bit.ly/ThreadPool>

# MySQL Enterprise Edition - SECURITY

## MySQL Enterprise TDE

- Data-at-Rest Encryption
- Key Management/Security

## MySQL Enterprise Encryption

- Public/Private Key Cryptography
- Asymmetric Encryption

## MySQL Enterprise Authentication

- External Authentication Modules
  - Microsoft AD, Linux PAMs, LDAP

## MySQL Data Masking

## MySQL Enterprise Firewall

- Block SQL Injection Attacks

## MySQL Enterprise Audit

## MySQL Enterprise Monitor

- Changes in Database Configurations, Users Permissions, Database Schema, Passwords

## MySQL Enterprise Backup

- Securing Backups, AES 256 encryption

## MySQL Enterprise Thread pool

- Attack Hardening

# Thank you

Contacts:

Catherine Schrimsher  
[catherine.schrimsher@oracle.com](mailto:catherine.schrimsher@oracle.com)

Eric Yanta  
[eric.yanta@oracle.com](mailto:eric.yanta@oracle.com)

---

